
Amazon Simple Storage Service

Getting Started Guide



Amazon Simple Storage Service: Getting Started Guide

Copyright © 2020 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

| | |
|--|----|
| Getting Started | 1 |
| Amazon S3 Basics | 1 |
| Sign Up for Amazon S3 | 2 |
| Create a Bucket | 3 |
| Add an Object to a Bucket | 5 |
| View an Object | 7 |
| Move an Object | 9 |
| Delete an Object and Bucket | 11 |
| Where Do I Go From Here? | 14 |
| Common Use Scenarios | 14 |
| Considerations Going Forward | 14 |
| AWS Account and Security Credentials | 15 |
| Security | 15 |
| AWS Integration | 15 |
| Pricing | 15 |
| Advanced Amazon S3 Features | 15 |
| Development Resources | 16 |
| Reference Resources | 16 |
| About This Guide | 17 |

Getting Started with Amazon Simple Storage Service

Amazon Simple Storage Service (Amazon S3) is storage for the internet. You can use Amazon S3 to store and retrieve any amount of data at any time, from anywhere on the web. You can accomplish these tasks using the AWS Management Console, which is a simple and intuitive web interface.

This guide introduces you to Amazon S3 and how to use the AWS Management Console to complete the following tasks:

- [Sign Up for Amazon S3 \(p. 2\)](#)
- [Create a Bucket \(p. 3\)](#)
- [Add an Object to a Bucket \(p. 5\)](#)
- [View an Object \(p. 7\)](#)
- [Move an Object \(p. 9\)](#)
- [Delete an Object and Bucket \(p. 11\)](#)

For information about Amazon S3 features, pricing, and frequently asked questions, see the [Amazon S3 product page](#).

Amazon S3 Basics

To get the most out of Amazon S3, you need to understand a few simple concepts. Amazon S3 stores data as objects within buckets. An object consists of a file and optionally any metadata that describes that file.

To store an object in Amazon S3, you upload the file you want to store to a bucket. When you upload a file, you can set permissions on the object and any metadata.

Buckets are the containers for objects. You can have one or more buckets. For each bucket, you can control access to it (who can create, delete, and list objects in the bucket), view access logs for it and its objects, and choose the geographical region where Amazon S3 will store the bucket and its contents.

Sign Up for Amazon S3

To use Amazon S3, you need an AWS account. If you don't already have one, you'll be prompted to create one when you sign up for Amazon S3. You are not charged for Amazon S3 until you use it. For information about Amazon S3 features, pricing, and frequently asked questions, see the [Amazon S3 product page](#).

To sign up for Amazon S3

1. Go to <https://aws.amazon.com/s3/> and choose **Get started with Amazon S3**.
2. Follow the on-screen instructions.

AWS notifies you by email when your account is active and available for you to use.

To start using Amazon S3, see [Create a Bucket \(p. 3\)](#).

Create a Bucket


Now that you've signed up for Amazon S3, you're ready to create a bucket using the AWS Management Console. Every object in Amazon S3 is stored in a bucket. Before you can store data in Amazon S3, you must create a bucket.

Note

You are not charged for creating a bucket. You are charged only for storing objects in the bucket and for transferring objects in and out of the bucket. The charges that you incur through following the examples in this guide are minimal (less than \$1). For more information about storage charges, see [Amazon S3 Pricing](#).

To create an S3 bucket

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.



3. In the **Bucket name** box, enter a unique DNS-compliant name for your new bucket.

The example screenshot uses the bucket name `admin-created`. You can't use this name because S3 bucket names must be unique. Create your own bucket name using the following naming guidelines:

- The name must be unique across all existing bucket names in Amazon S3.
- After you create the bucket, you cannot change the name, so choose wisely.
- Choose a bucket name that reflects the objects in the bucket because the bucket name is visible in the URL that points to the objects that you're going to put in your bucket.

For information about naming buckets, see [Rules for Bucket Naming](#) in the *Amazon Simple Storage Service Developer Guide*.

4. For **Region**, choose US West (Oregon) as the Region where you want the bucket to reside.
5. Choose **Create**.

The screenshot shows the 'Create bucket' wizard in the AWS Management Console. The wizard has a blue header with the title 'Create bucket' and a close button (X). Below the header is a progress bar with four steps: 1. Name and region, 2. Configure options, 3. Set permissions, and 4. Review. The first step, 'Name and region', is currently active. It contains a 'Bucket name' field with the text 'admin-created' and a 'Region' dropdown menu set to 'US West (Oregon)'. Below these fields is a section titled 'Copy settings from an existing bucket' with a dropdown menu that says 'Select bucket (optional)' and '47 Buckets'. At the bottom of the wizard, there are three buttons: 'Create' (highlighted with a red arrow), 'Cancel', and 'Next'.

You've created a bucket in Amazon S3.


To add an object to your bucket, see [Add an Object to a Bucket \(p. 5\)](#).

Add an Object to a Bucket


Now that you've created a bucket, you're ready to add an object to it. An object can be any kind of file: a text file, a photo, a video, and so on.

To upload an object to a bucket

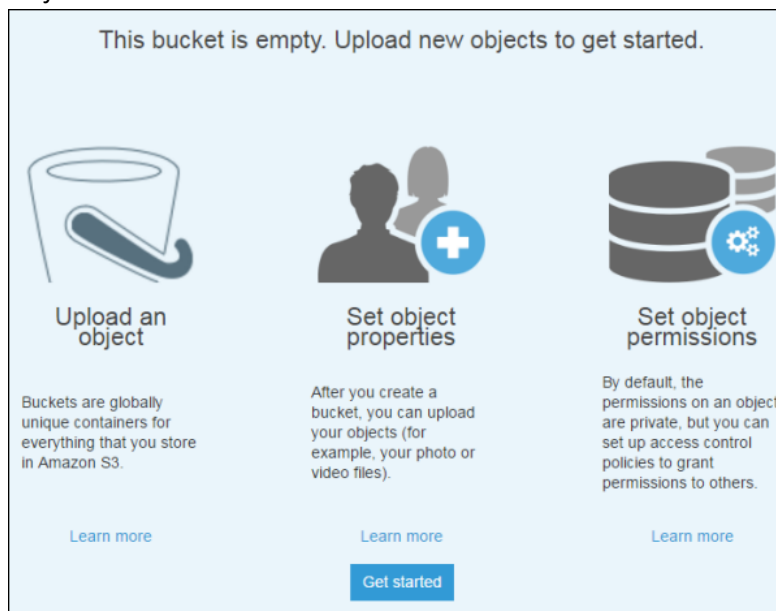
1. In the **Bucket name** list, choose the name of the bucket that you want to upload your object to.



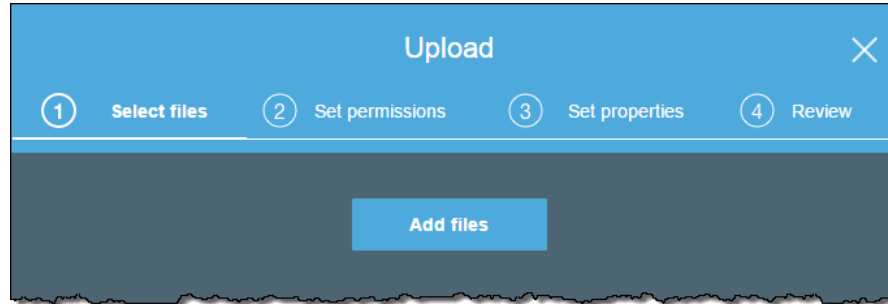
2. Choose **Upload**.




Or you can choose **Get started**.



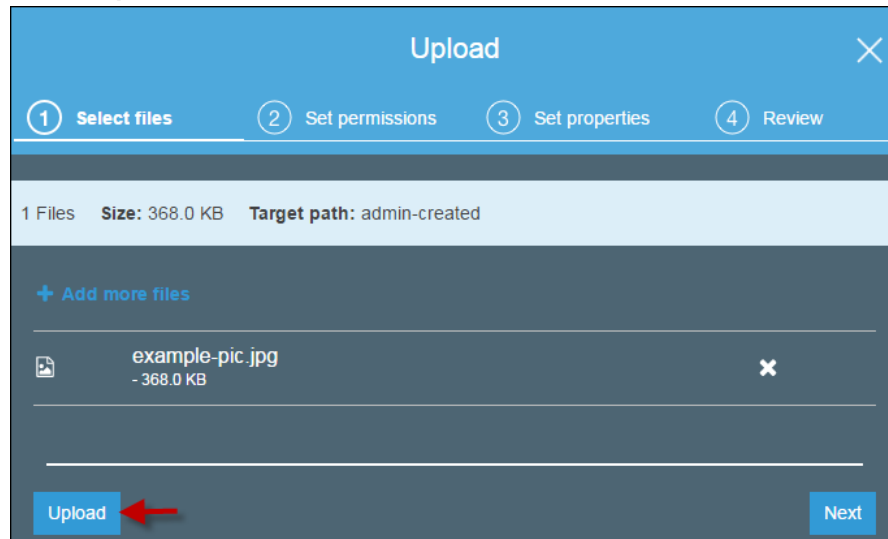
3. In the **Upload** dialog box, choose **Add files** to choose the file to upload.



4. Choose a file to upload, and then choose **Open**.



5. Choose **Upload**.



You've successfully added an object to your bucket.


To view your object, see [View an Object \(p. 7\)](#).

View an Object

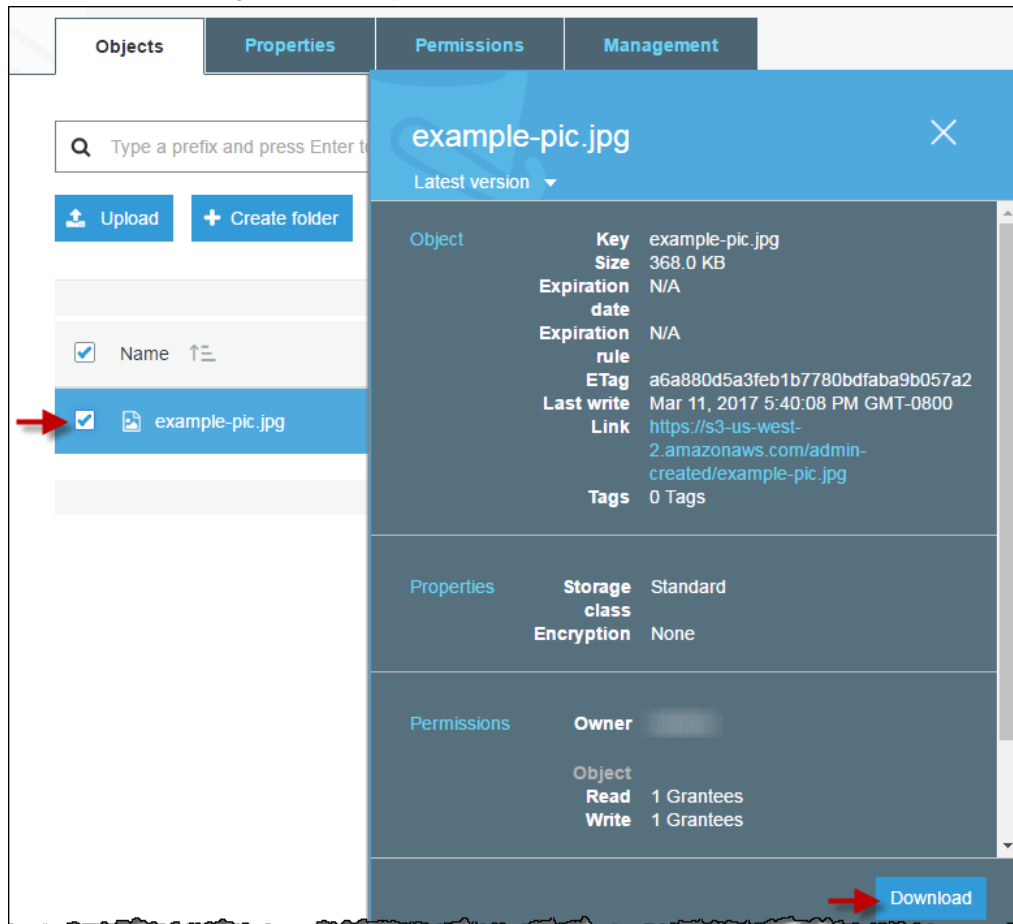
Now that you've added an object to a bucket, you can view information about your object and download the object to your local computer.

To download an object from a bucket

1. In the **Bucket name** list, choose the name of the bucket that you created.



2. In the **Name** list, select the check box next to the object that you uploaded, and then choose **Download** on the object overview panel.



You've successfully downloaded your object.

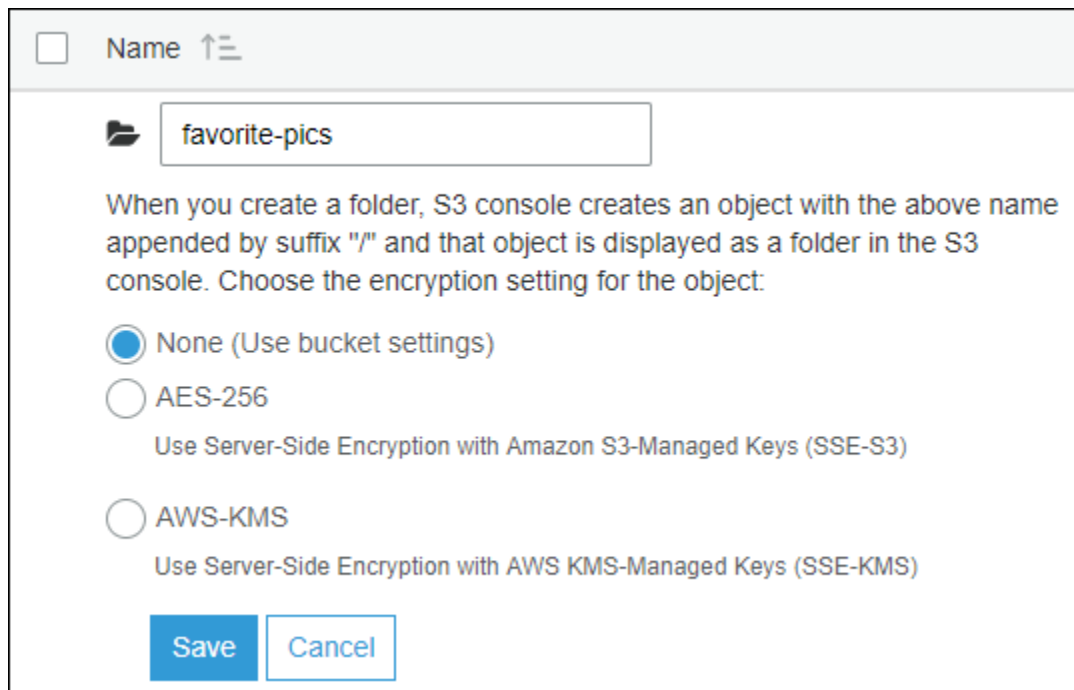
To copy and paste your object within Amazon S3, see [Move an Object \(p. 9\)](#).

Move an Object

So far you've added an object to a bucket and downloaded the object. Now you create a folder and move the object into the folder by copying and pasting the object.

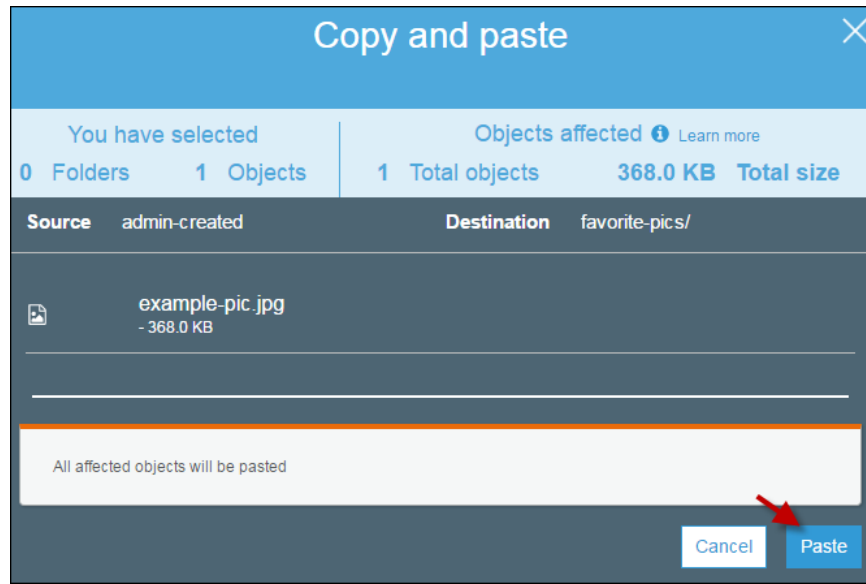
To copy an object

1. In the **Bucket name** list, choose the name of the bucket that you created.
2. Choose **Create Folder**.
3. Enter the information to create the folder:
 - a. Enter a folder name (for example, **favorite-pics**).
 - b. For the folder encryption setting, choose **None**.
 - c. Choose **Save**.



The screenshot shows the 'Create Folder' dialog box in the Amazon S3 console. At the top, there is a header bar with a checkbox and the text 'Name' followed by an upward arrow and a list icon. Below this, there is a folder icon and a text input field containing 'favorite-pics'. A paragraph of text explains that S3 creates an object with the folder name and a trailing slash, and that the object is displayed as a folder. Below this text are three radio button options for encryption: 'None (Use bucket settings)' (which is selected), 'AES-256', and 'AWS-KMS'. Under 'AES-256' and 'AWS-KMS' are sub-labels: 'Use Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3)' and 'Use Server-Side Encryption with AWS KMS-Managed Keys (SSE-KMS)' respectively. At the bottom are two buttons: 'Save' and 'Cancel'.

4. In the **Name** list, select the check box for the object that you want to copy.
5. Choose **Actions**, and then choose **Copy**.
6. In the **Name** list, choose the name of the folder, for example, **favorite-pics**.
7. Choose **Actions**, and then choose **Paste**.
8. Choose **Paste**.



The object appears in the folder where you pasted it.

To delete an object and a bucket in Amazon S3, see [Delete an Object and Bucket \(p. 11\)](#).

Delete an Object and Bucket


If you no longer need to store the object that you uploaded and made a copy of while going through this guide, you should delete the objects to prevent further charges.

You can delete the objects individually. Or you can empty a bucket, which deletes all the objects in the bucket without deleting the bucket.

You can also delete a bucket and all the objects contained in the bucket. However, if you want to continue to use the same bucket name, don't delete the bucket. We recommend that you empty the bucket and keep it. After a bucket is deleted, the name becomes available to reuse, but the name might not be available for you to reuse for various reasons. For example, it might take some time before the name can be reused, and some other account could create a bucket with that name before you do.

To delete an object from a bucket


1. In the **Bucket name** list, choose the name of the bucket that you want to delete an object from.
2. In the **Name** list, select the check box for the object that you want to delete.
3. Choose **Actions**, and then choose **Delete**.
4. In the **Delete objects** dialog box, verify that the name of the object, and then choose **Delete**.



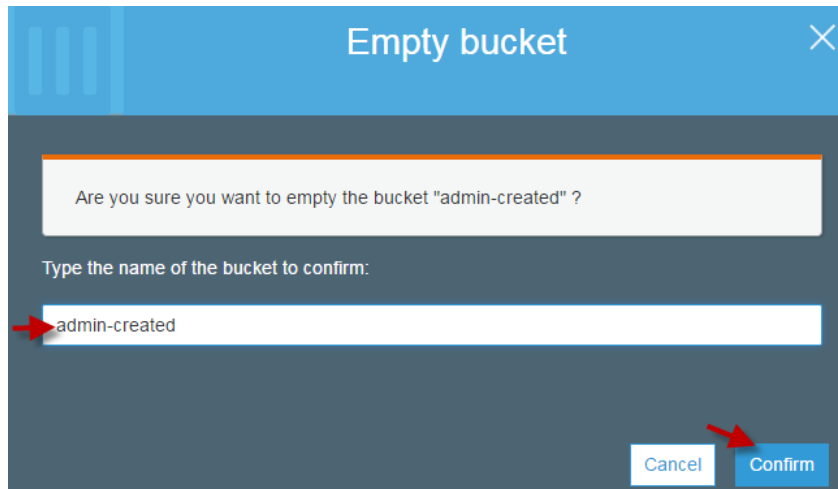
To empty a bucket

You can empty a bucket, which deletes all the objects in the bucket without deleting the bucket.

1. In the **Bucket name** list, choose the bucket icon next to the name of the bucket that you want to empty, and then choose **Empty bucket**.



2. To confirm emptying the bucket, in the **Empty bucket** dialog box, enter the name of the bucket, and choose **Confirm**.




To delete a bucket

You can delete a bucket and all the objects in it.


Important

If you want to continue to use the same bucket name, don't delete the bucket. We recommend that you empty the bucket and keep it. After a bucket is deleted, the name becomes available to reuse, but the name might not be available for you to reuse for various reasons.

1. In the **Bucket name** list, choose the bucket icon next to the name of the bucket that you want to delete, and then choose **Delete bucket**.



2. To confirm deletion, in the **Delete bucket** dialog box, enter the name of the bucket, and choose **Confirm**.



For more information about using Amazon S3, see [Where Do I Go From Here?](#) (p. 14).

Where Do I Go From Here?

In the preceding examples, you learned how to perform some basic Amazon S3 tasks. For more in-depth information, see one of the following Amazon S3 guides:

- The [Amazon Simple Storage Service Console User Guide](#) to learn more about using the Amazon S3 console.
- The [Amazon Simple Storage Service Developer Guide](#) to find detailed information about Amazon S3 features and code examples to support those features.
- The [Amazon Simple Storage Service API Reference](#) to find details about the Amazon S3 REST API.

The following topics explain various ways in which you can gain a deeper understanding of Amazon S3 so that you can implement it in your applications.

Topics

- [Common Use Scenarios](#) (p. 14)
- [Considerations Going Forward](#) (p. 14)
- [Advanced Amazon S3 Features](#) (p. 15)
- [Development Resources](#) (p. 16)
- [Reference Resources](#) (p. 16)

Common Use Scenarios

The AWS Solutions site lists many of the ways you can use Amazon S3. The following list summarizes some of those ways.

- **Backup and Storage** – Provide data backup and storage services for others.
- **Application Hosting** – Provide services that deploy, install, and manage web applications.
- **Media Hosting** – Build a redundant, scalable, and highly available infrastructure that hosts video, photo, or music uploads and downloads.
- **Software Delivery** – Host your software applications that customers can download.

For more information, see [AWS Solutions](#).

Considerations Going Forward

This section introduces you to topics you should consider before launching your own Amazon S3 product.

Topics

- [AWS Account and Security Credentials](#) (p. 15)
- [Security](#) (p. 15)
- [AWS Integration](#) (p. 15)
- [Pricing](#) (p. 15)

AWS Account and Security Credentials

When you signed up for the service, you created an AWS account using an email address and password. Those are your AWS account root user credentials. As a best practice, you should not use your root user credentials to access AWS. Nor should you give your credentials to anyone else. Instead, create individual users for those who need access to your AWS account. First, create an AWS Identity and Access Management (IAM) administrator user for yourself and use it for your daily work. For details, see [Creating Your First IAM Admin User and Group](#) in the *IAM User Guide*. Then create additional IAM users for other people. For details, see [Creating Your First IAM Delegated User and Group](#) in the *IAM User Guide*.

If you're an account owner or administrator and want to know more about IAM, see the product description at <https://aws.amazon.com/iam> or the technical documentation in the [IAM User Guide](#).

Security

Amazon S3 provides authentication mechanisms to secure data stored in Amazon S3 against unauthorized access. Unless you specify otherwise, only the AWS account owner can access data uploaded to Amazon S3. For more information about how to manage access to buckets and objects, go to [Identity and Access Management in Amazon S3](#) in the *Amazon Simple Storage Service Developer Guide*.

You can also encrypt your data before uploading it to Amazon S3.

AWS Integration

You can use Amazon S3 alone or in concert with one or more other Amazon products. The following are the most common products used with Amazon S3:

- [Amazon EC2](#)
- [Amazon EMR](#)
- [Amazon SQS](#)
- [Amazon CloudFront](#)

Pricing

Learn the pricing structure for storing and transferring data on Amazon S3. For more information, see [Amazon S3 pricing](#).

Advanced Amazon S3 Features

The examples in this guide show how to accomplish the basic tasks of creating a bucket, uploading and downloading data to and from it, and moving and deleting the data. The following table summarizes some of the most common advanced functionality offered by Amazon S3. Note that some advanced functionality is not available in the AWS Management Console and requires that you use the Amazon S3 API. All advanced functionality and how to use it is described in the [Amazon Simple Storage Service Developer Guide](#).

| Link | Functionality |
|--|--|
| Requester Pays Buckets | Learn how to configure a bucket so that a customer pays for the downloads they make. |

| Link | Functionality |
|---|---|
| Using BitTorrent With Amazon S3 | Use BitTorrent, which is an open, peer-to-peer protocol for distributing files. |
| Versioning | Learn about Amazon S3 versioning capabilities. |
| Hosting Static Websites | Learn how to host a static website on Amazon S3. |
| Object Lifecycle Management | Learn how to manage the lifecycle of objects in your bucket. Lifecycle management includes expiring objects and archiving objects (transitioning objects to the S3 <code>GLACIER</code> storage class). |

Development Resources

To help you build applications using the language of your choice, we provide the following resources:

- **Sample Code and Libraries** – The AWS Developer Center has sample code and libraries written especially for Amazon S3.

You can use these code samples as a means of understanding how to implement the Amazon S3 API. For more information, see the [AWS Developer Center](#).

- **Tutorials** – Our Resource Center offers more Amazon S3 tutorials.

These tutorials provide a hands-on approach for learning Amazon S3 functionality. For more information, see [Articles & Tutorials](#).

- **Customer Forum** – We recommend that you review the Amazon S3 forum to get an idea of what other users are doing and to benefit from the questions they ask.

The forum can help you understand what you can and can't do with Amazon S3. The forum also serves as a place for you to ask questions that other users or AWS representatives might answer. You can use the forum to report issues with the service or the API. For more information, see [Discussion Forums](#).

Reference Resources

The following list shows additional resources that you can use to further your understanding of Amazon S3.

- The [Amazon Simple Storage Service Console User Guide](#) describes all of the AWS Management Console functions related to Amazon S3.
- The [Amazon Simple Storage Service Developer Guide](#) provides a detailed discussion of the service.

It includes an architectural overview, detailed concept descriptions, and procedures for using the API.

- The [Amazon Simple Storage Service API Reference](#) provides a detailed discussion of the actions and parameters in Amazon S3.
- The Service Health Dashboard shows you the status of the Amazon S3 web service.

The dashboard shows you whether Amazon S3 (and all other AWS products) are functioning properly. For more information, see the [Service Health Dashboard](#).

About This Guide

This is the *Amazon Simple Storage Service Getting Started Guide*.

Amazon Simple Storage Service is frequently referred to within this guide as "Amazon S3." All copyrights and legal protections still apply.